

Natural Computation Methods in Machine Learning (NCML)

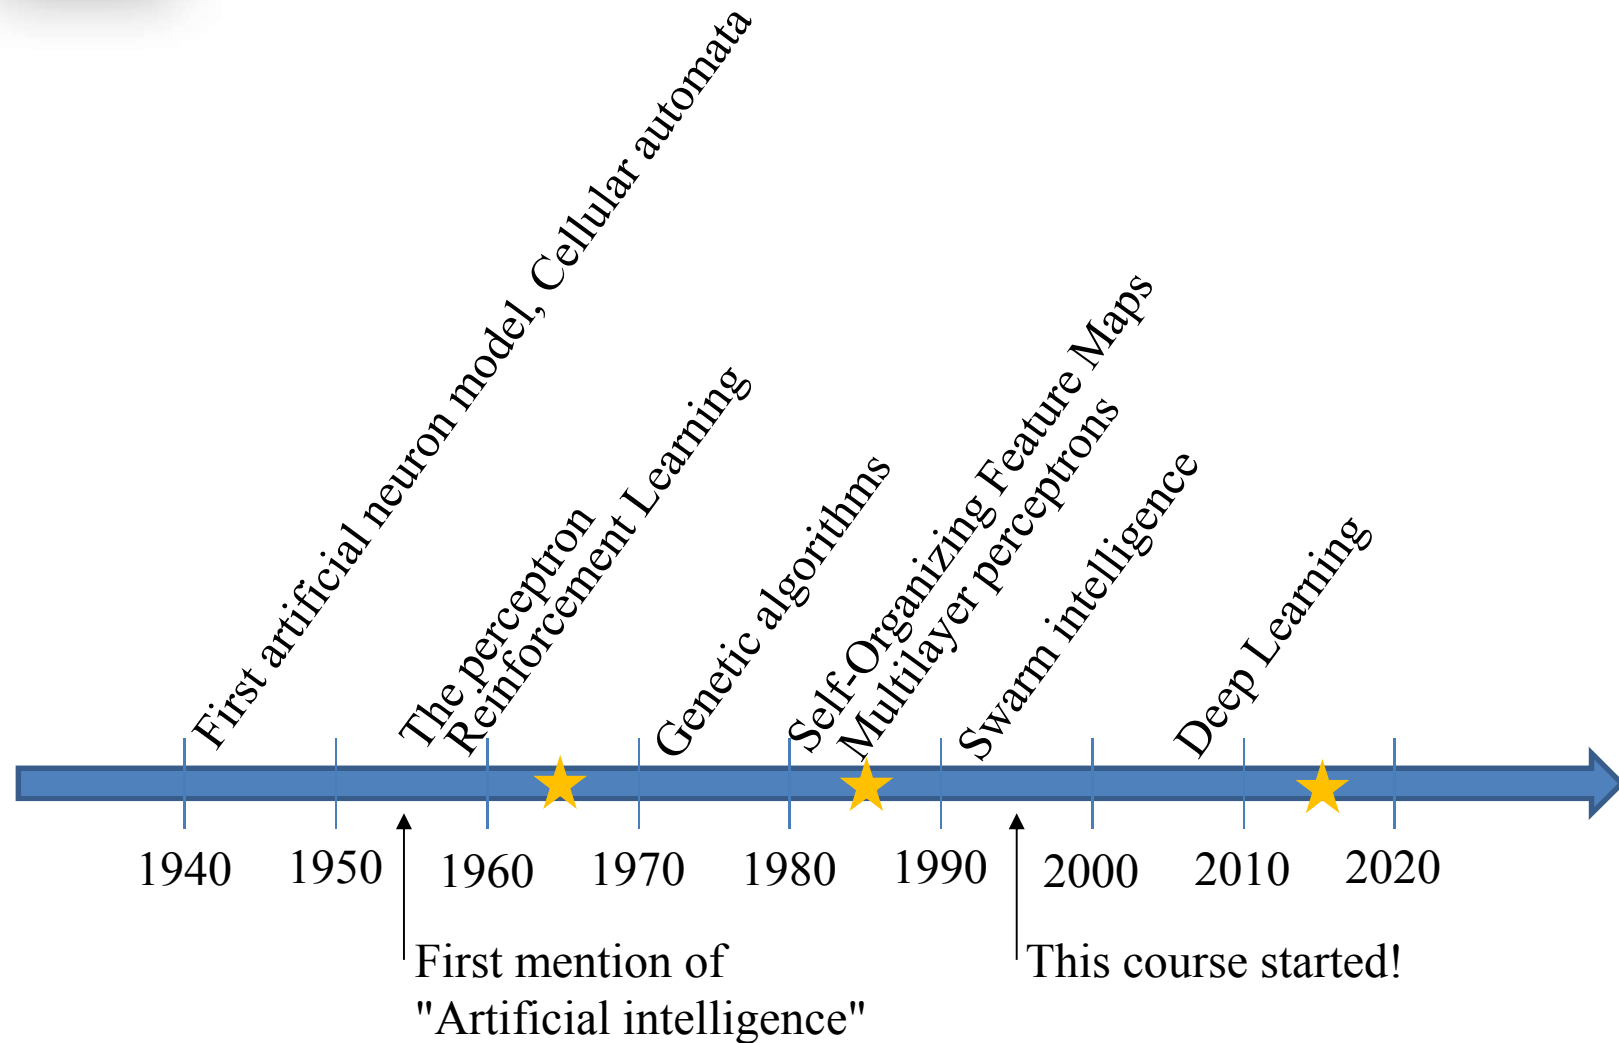
Lecture 2: Introduction to Artificial
Neural Networks (ANNs)

What is Artificial Intelligence?

- Modelling human intelligence?
 - No, not really (at least not in computer science)
- Automating tasks which we (previously) thought required intelligence
 - When done, this does not mean that the computer has become intelligent,
 - it more likely shows that the task did not require it
 - The solution is therefore, in hindsight, often no longer considered AI
 - Looking back, it therefore seems as if AI research never lead anywhere!
 - AI is a moving target!



Natural computation history



★ = AI/NC booms

Computers and humans



- Generalization
- Dealing with missing information
- Focusing on what is important
- Fault tolerance


Natural Computation

- If we want to solve problems for which we know nature/evolution has found solutions, it should be worth-while studying those solutions
- But don't take modelling too far!
 - It's a tool, not the goal
 - An airplane with flapping wings may be a better model (of birds), but that's not what aircraft engineers strive for. Usually ...






Artificial Neural Networks

- Interdisciplinary field
 - Many viewpoints
 - The same object may look very different from different angles (subjects)
 - Many different names for the same concepts
 - Sometimes difficult to search online for information (too many alternative keywords)

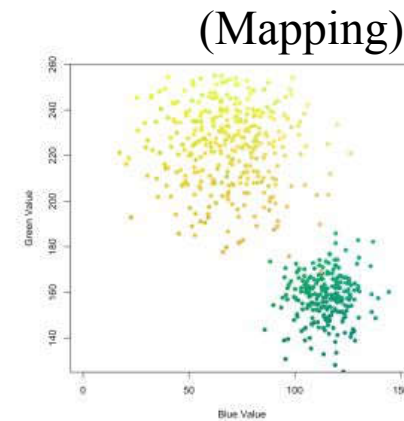



Neural networks

are universal function approximators



Olle (Classification)



$$\bar{y} = f(\bar{x})$$


Can approximate any function to any degree of accuracy

Not only for classification

- Some examples of applications which are not classification
 - Continuous control
 - Self-driving cars, for example, where the output (gas pedal and steering wheel) is continuous, even if there are many classification tasks to be solved as well
 - Value prediction
 - Stock market, weather, effects of climate change, probability that a loan applicant will repay the debt, ...
 - Extracting information from noisy data
 - Image analysis (not recognition – that's classification)
 - Compressing/decompressing, segmentation, enhancement, find and mark objects, ...

Neural networks



Biologically inspired, massively parallel networks of simple, adaptive, communicating nodes ('neurons')



$$\bar{y} = f(\bar{x})$$

(not necessarily 'layered' like this, but most often is)

Neurons from biology



The human body contains about 100 billion neurons
(\approx number of water drops required to fill an Olympic swimming pool)

The honeybee


© Totem Graphics Inc.



- 1 miljon brain cells in a few mm³
 - Approximately 10 TFLOP at 10 μ W
- With this computational device, it can*
- see, smell, recognize
 - walk, fly, maintain balance, navigate
 - communicate, find and collect food
 - ... and more
- Autonomously!*

The first artificial neuron model

McCulloch & Pitts 1943



$$y = f_h(S)$$

Activation function
(here a hard limiter)

$$f_h(S) = \begin{cases} 1, & \text{if } S > 0 \\ 0, & \text{if } S \leq 0 \end{cases}$$


$$S = \sum_{i=1}^n w_i x_i - \theta$$

threshold
or
bias

Very crude biological model,
but still useful for computations

Binary neurons = logic gates

Boolean logic



AND

a	b	y
0	0	0
0	1	0
1	0	0
1	1	1


$$y = f_h(S)$$

$$f_h(S) = \begin{cases} 1, & \text{if } S > 0 \\ 0, & \text{if } S \leq 0 \end{cases}$$

$$S = w_1 a + w_2 b - \theta$$

Binary neurons = logic gates

Boolean logic



AND

a	b	y
0	0	0
0	1	0
1	0	0
1	1	1


$$y = f_h(S)$$

$$f_h(S) = \begin{cases} 1, & \text{if } S > 0 \\ 0, & \text{if } S \leq 0 \end{cases}$$

$$S = w_1 a + w_2 b - \theta$$

Binary neurons = logic gates

Boolean logic



OR

a	b	y
0	0	0
0	1	1
1	0	1
1	1	1


$$y = f_h(S)$$

$$f_h(S) = \begin{cases} 1, & \text{if } S > 0 \\ 0, & \text{if } S \leq 0 \end{cases}$$

$$S = w_1 a + w_2 b - \theta$$

Binary neurons = logic gates

Boolean logic




$$y = f_h(S)$$

$$f_h(S) = \begin{cases} 1, & \text{if } S > 0 \\ 0, & \text{if } S \leq 0 \end{cases}$$

$$S = w_1 a + w_2 b - \theta$$

Binary neurons = logic gates

Boolean logic



NAND

a	b	y
0	0	1
0	1	1
1	0	1
1	1	0


$$y = f_h(S)$$

$$f_h(S) = \begin{cases} 1, & \text{if } S > 0 \\ 0, & \text{if } S \leq 0 \end{cases}$$

$$S = w_1 a + w_2 b - \theta$$

Binary neurons = logic gates

Boolean logic




NAND		
a	b	y
0	0	1
0	1	1
1	0	1
1	1	0

NAND is functionally complete – any Boolean function can be reformulated using NANDs only!

➔ A network of binary neurons is Turing complete

Challenge

- Set the weights of a binary neuron, so that it implements $(a \text{ AND } b) \text{ OR } c$



$$y = f_h(S)$$


$$f_h(S) = \begin{cases} 0, & \text{if } S \leq 0 \\ 1, & \text{if } S > 0 \end{cases}$$

$$S = w_1 a + w_2 b + w_3 c - \theta$$


a	b	c	y
0	0	0	0
0	0	1	1
0	1	0	0
0	1	1	1
1	0	0	0
1	0	1	1
1	1	0	1
1	1	1	1

Common properties of ANNs

- ANNs store information in the connections (weights), not in the nodes
- ANNs are fault tolerant
- ANNs are trained (by modifying the weights), not programmed
- ANNs can generalize – work in situations slightly different than before (without retraining)
- ANNs are adaptive – can be retrained if generalization does not suffice
- ANNs are concurrent



Learning




- Hebb's rule (1949)
 - Memory model for (biological) neural networks
 - If two nodes are active simultaneously, reinforce the connection between them
- Rosenblatt's Perceptron Convergence Procedure (PCP, 1958)
 - Started the 1960's boom
 - Next lecture

Three forms of learning

(you tried all three in the intro-lab)


Supervised



Imitation

PCP, Backprop,
Rprop, ...


Reinforcement



Trial-and-error

TD(λ), Sarsa,
Q-Learning, ...

Unsupervised




Structure

Hebb, SCL,
K-Means, SOFM

Network architectures


Feed-forward – information flows in one direction



- Used for classification, function approximation (regression), perception
- e.g. Multilayer perceptrons, CNNs

Network architectures

Recurrent – layered networks with loops





- Used for recognizing/generating sequences
- e.g. Jordan networks, Elman networks, LSTMs
(LSTM layers are very different though)

Network architectures

Fully interconnected recurrent

- Finite state machine!
- Used as associative memories, or for combinatorial optimization problems
- Training – often some extension of Hebb's rule
- The first neurocomputer, SNARC (1951) was based on this



A SNARC neuron


Why neural networks?

- Why not use statistical methods?
- Neural networks are statistical methods! (not as "model free" as sometimes claimed)
- Today, neural networks excel in many application areas, but they have been used for a long time for other reasons as well:
 - Rapid response (feed-forward networks)
 - hardware
 - trade-off training/recall
 - Rapid prototyping, e.g. NetTalk

NetTalk

Sejnowski & Rosenberg, 1987

- Learns to pronounce written text
 - Produces phonemes



- Did not beat state-of-the-art at the time, but close, and at a very small fraction of R&D time

ANN application classics

Tool boxes, adaptive filters, OCR, Image analysis, signal processing, touch pads, fraud detection, recommendation systems, ...

