

Natural Computation Methods in Machine Learning (NCML)

Lecture 3: The Perceptron


Pattern recognition

- Pattern recognition = Feature extraction + Classification
- Feature extraction: To find 'good' features
→ feature vector, \bar{x}
- Very sensitive to assumptions!
 - What's the most significant features to extract, to recognize someone's face?
 - prominent \neq significant
 - Lesson's learned from a masquerade ball




Classification

- Example: Measure two features
 - 2D feature vector, $\bar{x} = (x_1, x_2)$
- Classification: To find a *discriminant* that separates the classes, in the feature space (in this case 2D)
 - In the example, a line would suffice
 - Discriminants can take any shape
 - There is usually an infinite number of solutions (possible discriminants), even if restricted to a certain shape



Nearest neighbour classifiers

- Classify the unknown sample (vector/pattern) to the class of its closest previously classified neighbour



- Problem: The closest neighbour may be an outlier from the other class
- Solution: *K*-nearest-neighbour (KNN) – classify \bar{x} to the most common class among its *K* closest neighbours
- *Q: No explicit mention of finding discriminants, but there must be one. What's its shape?*
- *Q: Is this supervised or unsupervised learning?*


Distance measures

- Define distance between two vectors
 - $\bar{a} = [a_1, a_2, \dots, a_n]$ and $\bar{b} = [b_1, b_2, \dots, b_n]$
- General: the l_p norm of a vector, $\bar{x} = |\bar{a} - \bar{b}|$

$$l_p(\bar{x}) = \left(\sum_i x_i^p \right)^{\frac{1}{p}}$$

			\bar{b}
\bar{a}			

- Two well known special cases
 - $p=2 \rightarrow l_2$ norm = Euclidean distance
 - $p=1 \rightarrow l_1$ norm = city block (Manhattan) distance



l_p norm examples

$$l_p(\bar{x}) = \left(\sum_i x_i^p \right)^{\frac{1}{p}}$$


- Euclidean distance (l_2 norm)

$$d_{euc}(\bar{a}, \bar{b}) = l_2(|\bar{a} - \bar{b}|) = \sqrt{\sum_i |a_i - b_i|^2} = \sqrt{3^2 + 4^2} = \sqrt{25} = 5$$

- City block (Manhattan) distance (l_1 norm)



$$d_{cb}(\bar{a}, \bar{b}) = l_1(|\bar{a} - \bar{b}|) = \sum_i |a_i - b_i| = 3 + 4 = 7$$

- Affects shape of equi-distant border



The (binary) Perceptron

Rosenblatt 1958



$$y = f_h(S) = \begin{cases} 1, & \text{if } S > 0 \\ 0, & \text{if } S \leq 0 \end{cases}$$

$$S = \sum_{i=1}^n w_i x_i - \theta = \underbrace{\sum_{i=0}^n w_i x_i}_{\text{Augmented vector notation}}, \text{ where } \begin{cases} x_0 = -1 \\ w_0 = \theta \end{cases}$$

Augmented vector notation

Q: Here, $x_0 = -1$, but this would work for any value of x_0 , positive or negative. Why?

Hyperplanes


- The perceptron defines a hyperplane!
- 2D (hyperplane=line) example:

$$S = \sum_{i=1}^2 w_i x_i - \theta = w_1 x_1 + w_2 x_2 - \theta$$

- The perceptron 'flips' at $S=0$
 - so, set $S=0$ and solve for x_2

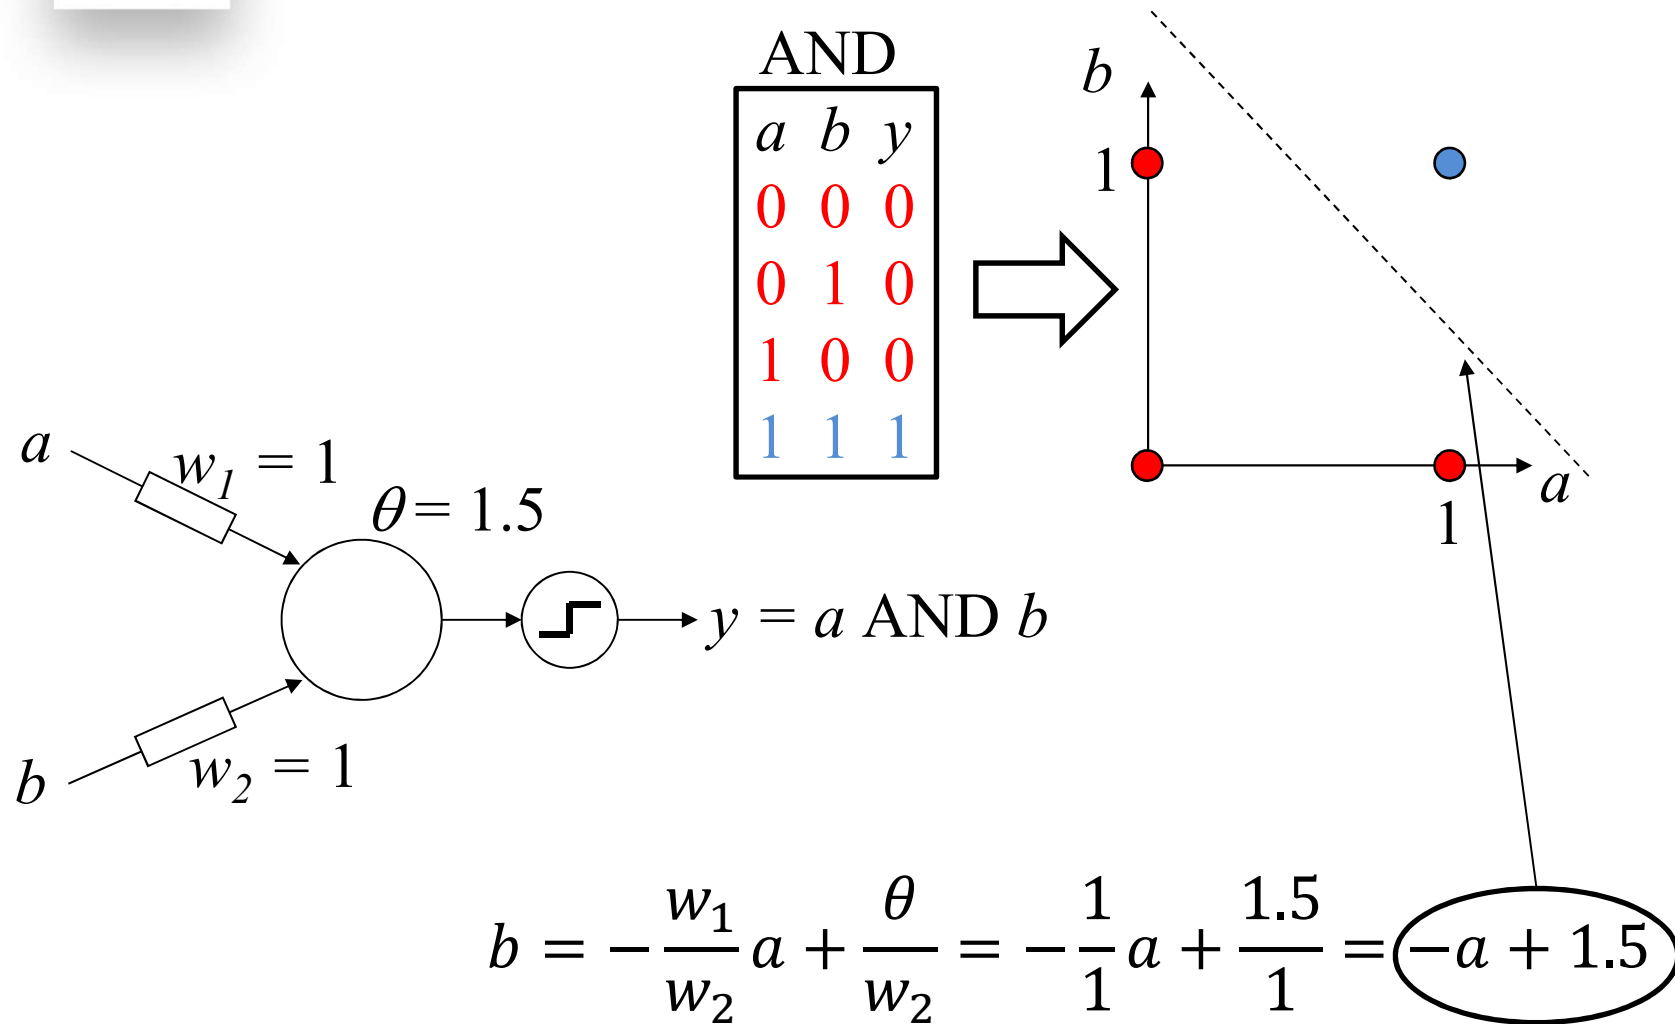
$$w_1 x_1 + w_2 x_2 - \theta = 0$$

$$x_2 = \frac{\theta - w_1 x_1}{w_2} = -\frac{w_1}{w_2} x_1 + \frac{\theta}{w_2} \quad (= kx+m \dots a \text{ line!})$$



\therefore The weights define the position and slope of a hyper plane!

Remember the AND gate?





Q: What about NAND?

What about the OR?

OR

a	b	y
0	0	0
0	1	1
1	0	1
1	1	1



$$b = -\frac{w_1}{w_2}a + \frac{\theta}{w_2} = -\frac{1}{1}a + \frac{0.5}{1} = -a + 0.5$$

Adjusting the line automatically

Training

- We need a number of pairs (\bar{x}, d) of feature vectors (\bar{x}) and desired responses (d)
- For each such pair, and perceptron output $y(\bar{x})$:
 - If $y=d$, do nothing
 - If $y=0, d=1$: Reinforce the connections (to increase S)
 - If $y=1, d=0$: Weaken the connections (to decrease S)
- Reinforce/weaken = add/subtract x_i to/from w_i
 - note that this will leave the weight unchanged if $x_i=0$
(makes sense since 0 inputs do not contribute to the error)

Perceptron Convergence Procedure

PCP, Rosenblatt 1958

w_i , $0 \leq i \leq n$, the weight from input i

x_i , $0 \leq i \leq n$, the value of input i

$x_0 = -1$ (or any other constant value)

- 1) Initialize the node: Set all weights, w_i , to small random values
- 2) Present a feature vector $\bar{x} = [x_1, x_2, \dots, x_n]$ and a desired response, d
- 3) Compute perceptron output: $y = f_h(S) = \begin{cases} 1, & S > 0 \\ 0, & S \leq 0 \end{cases} \quad S = \sum_{i=0}^n w_i x_i$
- 4) Adjust weights:

$$w_i(t+1) = w_i(t) + \Delta w_i$$

$$\Delta w_i = \begin{cases} 0, & \text{if } y=d \\ +x_i, & \text{if } y=0, d=1 \\ -x_i, & \text{if } y=1, d=0 \end{cases}$$

- 5) Repeat from 2) with a new vector

Perceptron Convergence Procedure

slightly rewritten, now with a throttle/brake (η)

w_i , $0 \leq i \leq n$, the weight from input i

x_i , $0 \leq i \leq n$, the value of input i

$x_0 = -1$ (or any other constant value)

- 1) Initialize the node: Set all weights, w_i , to small random values
- 2) Present a feature vector $\bar{x} = [x_1, x_2, \dots, x_n]$ and a desired response, d
- 3) Compute perceptron output: $y = f_h(S) = \begin{cases} 1, & S > 0 \\ 0, & S \leq 0 \end{cases} \quad S = \sum_{i=0}^n w_i x_i$
- 4) Adjust weights:

$$w_i(t+1) = w_i(t) + \Delta w_i$$

$\Delta w_i = \eta \delta x_i$

, where $\delta = d - y$




Many update rules can be written like this, only the definition of δ differs.

- 5) Repeat from 2) with a new vector

Perceptron Convergence Theorem

PCT

- The Perceptron Convergence Procedure converges to an optimal discriminant in a finite number of steps, if such a discriminant exists
- PCP + PCT → The boom in the 1960's



Mark I (1957): One the first neurocomputers (the first based on Perceptrons)

Adam


The first European digital neurocomputer



Perceptron Convergence Theorem

PCT

- The Perceptron Convergence Procedure converges to an optimal discriminant in a finite number of steps, ***if such a discriminant exists***
 - Problem: It seldom does! Few interesting classification problems are separable by one linear discriminant (a hyperplane)
 - Simple example: XOR



Sudden death of the first boom


- Minsky & Papert, *Perceptrons*, 1969
- Old news to the researchers, but a shock to the funding agencies ...
- No grants → Mass exodus to other fields
 - Bernard Widrow, for example, removed the main problem (the step function) and thus invented the adaptive filter
- Neural winter, 1969 – early 1980's



What can we do?

Solution 1: Linearize the problem

- Increase dimensionality – add more inputs
 - For example, XOR becomes linearly separable if we add an extra input c which is ' a AND b '



XOR(a, b)

a	b	c	y
0	0	0	0
0	1	0	1
1	0	0	1
1	1	1	0


- Important (and often used) general principle:
Complex problems may become easier by first projecting them to a higher-dimensional space

Challenge: Geometrical explanation why this makes XOR linearly separable

What can we do?

Solution 1: Linearize the problem

- Increase dimensionality – add more inputs
 - For example, XOR becomes linearly separable if we add an extra input c which is ' a AND b '




XOR(a, b)

a	b	c	y
0	0	0	0
0	1	0	1
1	0	0	1
1	1	1	0

- Important (and often used) general principle:
Complex problems may become easier by first projecting them to a higher-dimensional space


Challenge: Geometrical explanation why this makes XOR linearly separable

OR 2D→3D



What can we do?

Solution 2: Combine perceptrons in layers



But how do we train such a network? (next lecture)