


Natural Computation Methods in Machine Learning (NCML)

Lecture 4: Multilayer Perceptrons and
Backpropagation





From previous lecture

Linear (in)separability

- The Perceptron Convergence Procedure converges to an optimal discriminant in a finite number of steps, ***if such a discriminant exists***
 - Problem: It seldom does! Few interesting classification problems are separable by one linear discriminant (a hyperplane)
 - Simple example: XOR




From previous lecture



From previous lecture

Multilayer perceptrons (MLP)




Problem: How to find the weight values automatically
(i.e. how to train the network)

The credit assignment problem

Structural. We will discuss a temporal one later


- To decide how much to blame an individual weight for the result
- We now have '*hidden layers*'
 - No target (desired output) information
 - What is the error of a hidden node?
- The step function (f_h) is an obstacle
 - Can't decide, analytically from the outputs, how close to the flipping point ($S=0$) the weighted sum was
 - Removing f_h ($y=S$) does not help (*us. It helped B. Widrow*)
 - Linear nodes \rightarrow The MLP can be reduced to one (output) layer
 - Back to where we started



\therefore The activation function, $f(S)$, must be non-linear and differentiable

Sigmoid functions

- Sigmoid = any S-shaped function
- Often confused with the *logistic function*, which is actually just an example, though a very common one:



- $\lambda (\geq 0)$ decides slope. In the extremes:
 - $\lambda = 0 \rightarrow y = 0.5$ (flat line)
 - $\lambda \rightarrow \infty \rightarrow y = f_h(S)$ (step function)
- Simple derivative: $y' = f'(S) = \lambda y(1 - y)$
- Another commonly used sigmoid function: $\tanh(S)$
 \approx logistic, but in the range $] -1, 1[$ instead of $] 0, 1[$

How to make a learning rule

A (very) general recipe


1. State what is to be minimized as a *loss function*^(*)

For example:

$$E = \frac{1}{2} (d - y)^2 \quad \begin{array}{l} d = \text{desired output} \\ y = \text{actual output} \end{array}$$

2. How much did weight w_i , contribute to this loss?

$$\frac{\partial E}{\partial w_i}$$



3. Make an update rule which moves the weight in proportion to its contribution, but in the other direction:

$$w_i \leftarrow w_i + \Delta w_i \quad \Delta w_i = -\eta \frac{\partial E}{\partial w_i}$$

4. Can often be expressed as $\Delta w_i = \eta \delta x_i$ where x_i is the input corresponding to weight w_i

^(*)*loss function = error function = cost function \approx objective function*


Deriving the delta rule

following the recipe

- Consider a single neuron with a differentiable activation function:

$$y = f(S)$$

$$S = \sum_{i=0}^n w_i x_i, \text{ where } \begin{cases} x_0 = -1 \\ w_0 = \theta \end{cases}$$



- Step 1: State what is to be minimized*
 - Let's assume that the loss function is the squared error

$$E = \frac{1}{2}(d - y)^2 \quad (\text{hidden assumption here – 'Gaussian prior'. We assume that data comes from a normal distribution})$$

- Step 2: Use the chain rule to break down $\frac{\partial E}{\partial w_i}$*

Deriving the delta rule

Use the chain rule to break down $\partial E / \partial w_i$

- The loss (E) depends on the output (y), which depends on the sum (S), which depends on the weight (w_i):


$$\frac{\partial E}{\partial w_i} = \frac{\partial E}{\partial y} \frac{\partial y}{\partial S} \frac{\partial S}{\partial w_i}$$

$$\frac{\partial E}{\partial y} = \frac{\partial \frac{1}{2}(d - y)^2}{\partial y} = 2 \frac{1}{2}(d - y)(-1) = -(d - y)$$

$$\frac{\partial y}{\partial S} = \frac{\partial f(S)}{\partial S} = f'(S)$$

$$\frac{\partial S}{\partial w_i} = \frac{\partial \sum_{j=0}^n w_j x_j}{\partial w_i} = \frac{\partial (w_i x_i)}{\partial w_i} = x_i$$

$$\therefore \frac{\partial E}{\partial w_i} = -(d - y)f'(S)x_i$$



The delta rule

Express Δw

- *Step 3:* $\Delta w_i = -\eta \frac{\partial E}{\partial w_i}$ $\frac{\partial E}{\partial w_i} = -(d - y)f'(S)x_i$

$$\Delta w_i = \eta \underbrace{(d - y)f'(S)}_{\delta} x_i$$

- Or, in the general form:


$$\Delta w_i = \eta \delta x_i \quad \text{where } \delta = f'(S)(d - y)$$

- If we assume that the activation function is logistic:

$$\delta = \lambda y(1 - y)(d - y)$$

- This is the **delta rule**, a.k.a. LMS (least-mean-squares) (Widrow & Hoff, 1960)


- Compare to the Perceptron Convergence Procedure (PCP)!
- Same. but without $f'(S)$, since the step function used in a binary perceptron is not differentiable



Backpropagation

The generalized delta rule

- We can extend the delta rule to cover a whole network of neurons (as long as everything is differentiable)
 - Same idea, the chain of partial derivatives just gets longer
- But, we now have several nodes. Therefore:
 - We must index the nodes, δ -values, and desired outputs d
 - Weights need a second index. Let w_{ji} denote the weight from node i to node j
 - x_i in the equations is the input to the node, i.e. it may be the value of a hidden node, not necessarily an input to the network.



Backpropagation

"of errors"

- We can still express the update rule on the general form

$$\Delta w_{ji} = \eta \delta_j x_i$$


- but the definition of δ_j now depends on if node j is an output node or a hidden node
- For an output node, the delta rule is applicable as it is:

$$\delta_j = f'(S_j)(d_j - y_j) = \underbrace{\lambda y_j(1 - y_j)}_{\substack{\uparrow \\ \text{(if logistic)}}}(d_j - y_j)$$

Backpropagation

"of errors"

- Let's assume, for now, that we only have one hidden layer
- For a hidden node there is no desired output, d_j , but:
 - The hidden layer contributes to the output error through the *hidden-to-output* weights (through weighted sums)
 - The hidden layer should therefore be blamed for the error, in proportion to those same weights
 - The error of a hidden node is a weighted sum of the δ -values we just computed for the outputs
 - In other words, we backpropagate errors



$$\delta_j = f'(S_j) \sum_k w_{kj} \delta_k = \lambda y_j (1 - y_j) \sum_k w_{kj} \delta_k$$

where the sum is over the nodes in the next layer

Backpropagation Algorithm

1. Initialize. Set all weights to small random values with zero mean
2. Present an input vector, $\bar{x} = (x_1, x_2, \dots, x_n)$, and corresponding target vector, $\bar{d} = (d_1, d_2, \dots, d_m)$
3. Feed forward phase (*recall*): Compute network outputs, by updating the nodes layer by layer from the first hidden layer to the outputs. The first hidden layer computes (for all nodes, y_j):

$$y_j = f \left(\sum_{i=0}^n w_{ji} x_i \right), \text{ where } x_0 = -1$$

The next layer applies the same formula, substituting this layer's node values for x_i , etc.

Back propagation Algorithm

4. Back propagation phase: Compute weight changes^(*) iteratively, layer by layer, from the outputs to the first hidden layer:

$$\Delta w_{ji} = \eta \delta_j x_i$$

$$\delta_j = \begin{cases} \lambda y_j (1 - y_j) (d_j - y_j), & \text{if } y_j \text{ is an output node} \\ \lambda y_j (1 - y_j) \sum_k w_{kj} \delta_k, & \text{if } y_j \text{ is a hidden node} \end{cases}$$

(The sum is over all k nodes in the next layer i.e. the layer for which δ -values were computed in the previous iteration)

5. Repeat from step 2 with a new input-target pair

() Note that we only compute weight changes here. It does not say when to update the weight.*

Back propagation

Some notes

- λ is redundant here. It can be embedded in η .
 - λ is therefore often ignored (assumed to be 1)
- Weighted sums are just matrix-vector products.

When computing the output:


$$y_j = f\left(\sum_{i=0}^n w_{ji}x_i\right) = f(W\bar{x})$$

- so if you implement this using a matrix library, you can compute the outputs of all nodes in a layer in one shot
- If so, when backpropagating δ -values, just transpose the matrix!

$$\sum_k w_{kj}\delta_k = W^T \bar{\delta}$$

Gradient descent

- Backprop implements *gradient descent* (Cauchy, 1847)



- The delta rule was invented in 1960 (by Widrow & Hoff)
- Why didn't Minsky&Papert see this solution in 1969?
 - Can get stuck in closest local minimum (almost certainly will)
 - Would have been considered a big problem in the 1960's
 - The cost of re-starting an experiment was very high

Momentum

- Common improvement: Add a *momentum term* to the weight update

$$\Delta w_{ji}(t + 1) = \eta \delta_j x_i + \alpha \Delta w_{ji}(t)$$

- Smoothing out weight changes over time
 - Gives the 'ball' a momentum, i.e. tends to continue in the direction as before
- Similar effects if we adapt step length over time
 - next lecture

When to update the weights

Epoch learning v.s. Pattern learning

- Epoch learning
 - Accumulate Δw until all patterns^(*) have been presented once (= 1 epoch). Then update the weight and clear Δw
 - Special case of Batch Learning (batches can be smaller than the whole training set)
- Pattern learning (stochastic)
 - Update w after each pattern presentation
 - as a new step 4.5: $w_{ji} \leftarrow w_{ji} + \Delta w_{ji}$
 - Requires random order of presentation (hence 'stochastic')
 - In this special case = *stochastic gradient descent*

^(*) *Patterns = Input vectors*

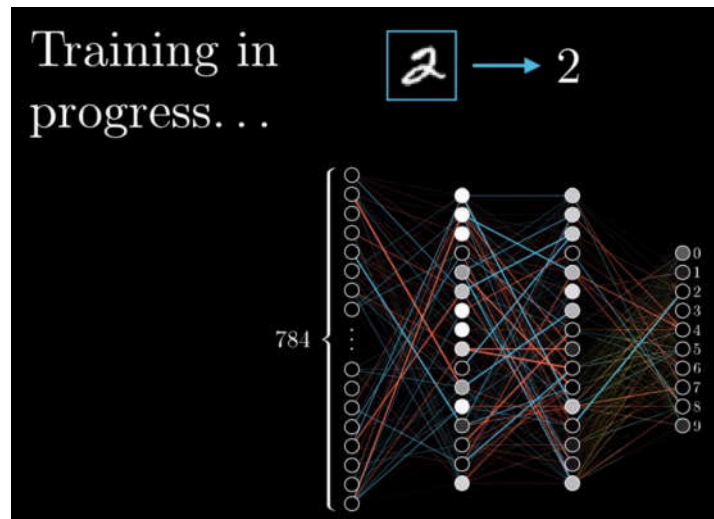
When to update the weights


Epoch learning v.s. Pattern learning

- Epoch learning
 - This is what makes Backprop = gradient descent
 - Theorems and algorithm variants, often require this
- Pattern learning (stochastic)
 - Often better in practice (if the algorithm allows is)
 - The non-determinism reduces risk of getting stuck
 - Usually converges faster
- Common compromise: Batch learning, for smaller subsets – "mini-batches"

Backprop videos on YouTube

- There are many, of course
- One of my favourite channels: 3Blue1Brown
 - For example, watch the video "What is backpropagation really doing?" (from 3:09)
 - (<https://www.youtube.com/watch?v=Ilg3gGewQ5U>)





Challenges

All are possible exam questions

- The loss function could be any differentiable function, but is very often assumed to be the squared error
 - Which part of the update equations would change, if we replaced the objective function?
- If we use pattern learning, why don't we update the weight directly, in step 4? (instead of as a new step 4.5)
- What would happen if we initialized all weights to zero, instead of small random values?
- In a deep network (many layers) the chain of partial derivatives gets very long. A problem?

Super Challenge

- As shown in the previous lecture and this one:
 $ab, \overline{ab}, a + b, \overline{a + b}$, are all linearly separable,
 $(a + b)\overline{c}$ (the three input solution to XOR) is too,
but $a \oplus b = a\overline{b} + \overline{a}b = (a + b)\overline{ab}$ is not!
- *Where is the limit? Under which condition(s) is a Boolean expression linearly separable, when viewed as a classification problem?*

*(This challenge is beyond the scope of this course,
so don't worry if you can't figure this one out)*