

Natural Computation Methods in Machine Learning (NCML)

Lecture 10: Growing Neural Gas
(and a recap of CL and SOFM)


CL and SOFM, revisited

- A node's *position* (in the input space) is its weight vector
 - Nodes move around in the same space as the data
- The purpose of unsupervised learning is (usually) to cluster for classification, or for modelling distributions
 - We want several nodes in high density areas, fewer in low density areas
 - i.e. we want node distributions which follow the data distributions

Competitive Learning

Effects of bad initialization

- Bad utilization if we randomize initial positions

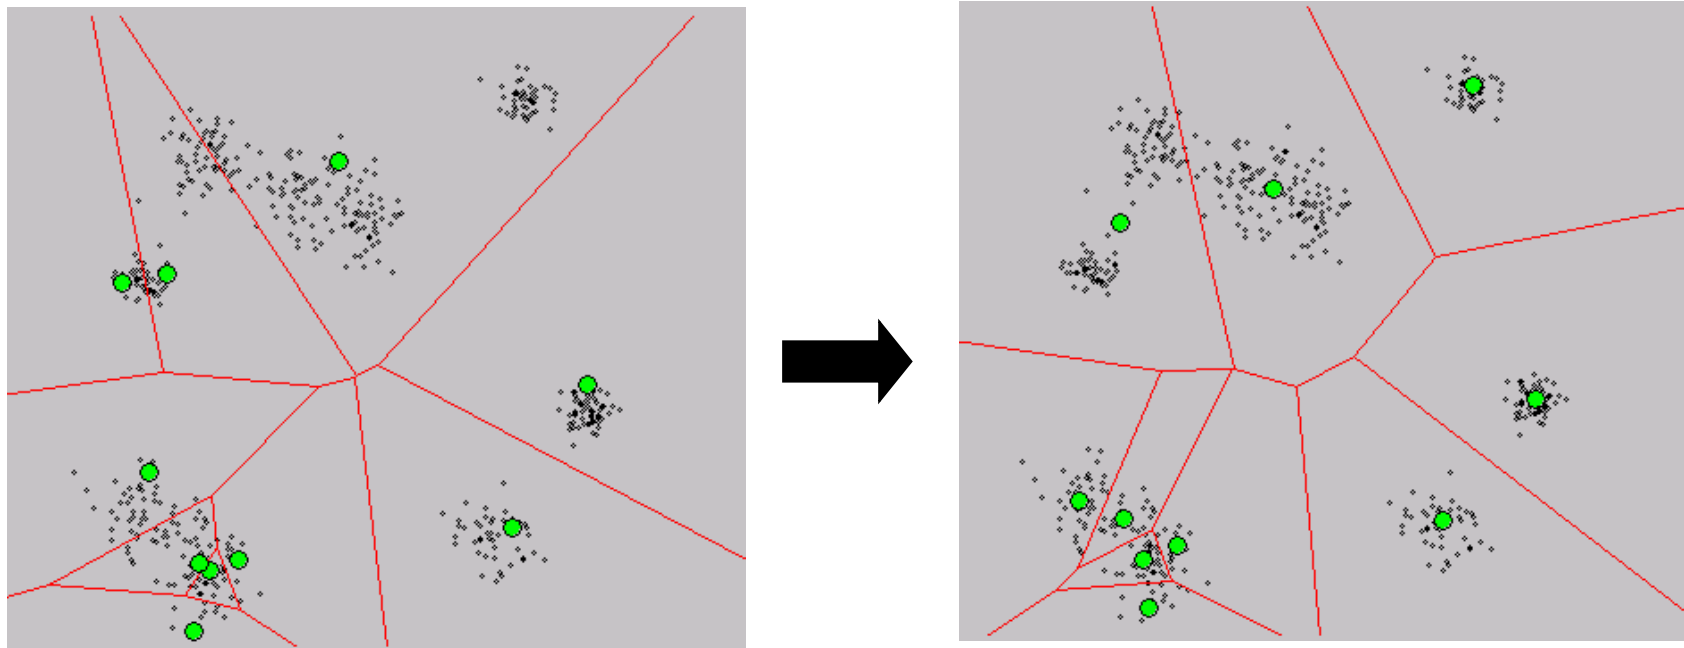


*Some nodes (green) are unused – they will never move
since their Voronoi regions are empty*

Competitive Learning

Effects of bad initialization

- Better if initialized from the data instead
- Still not perfect though (unfair distribution of nodes)
 - Some nodes seem to need help to cover their data clusters



Q: How can we detect, automatically, that a node 'needs help'?


A: It will move around more than the others

Competitive Learning

Problems with non-stationary distributions

- What should we expect if we run competitive learning with 10 nodes on data from this uniform distribution?

- They should spread out approximately uniformly




- What if the square moves, very slowly?
 - Works only if it moves very slowly, leaving no nodes behind (without any data in their Voronoi regions)
- What if it moves faster, or jumps?
 - The Winner-takes-all scenario strikes again!
- What if this had been a Self-Organizing Feature Map?

Self-Organizing Feature Maps

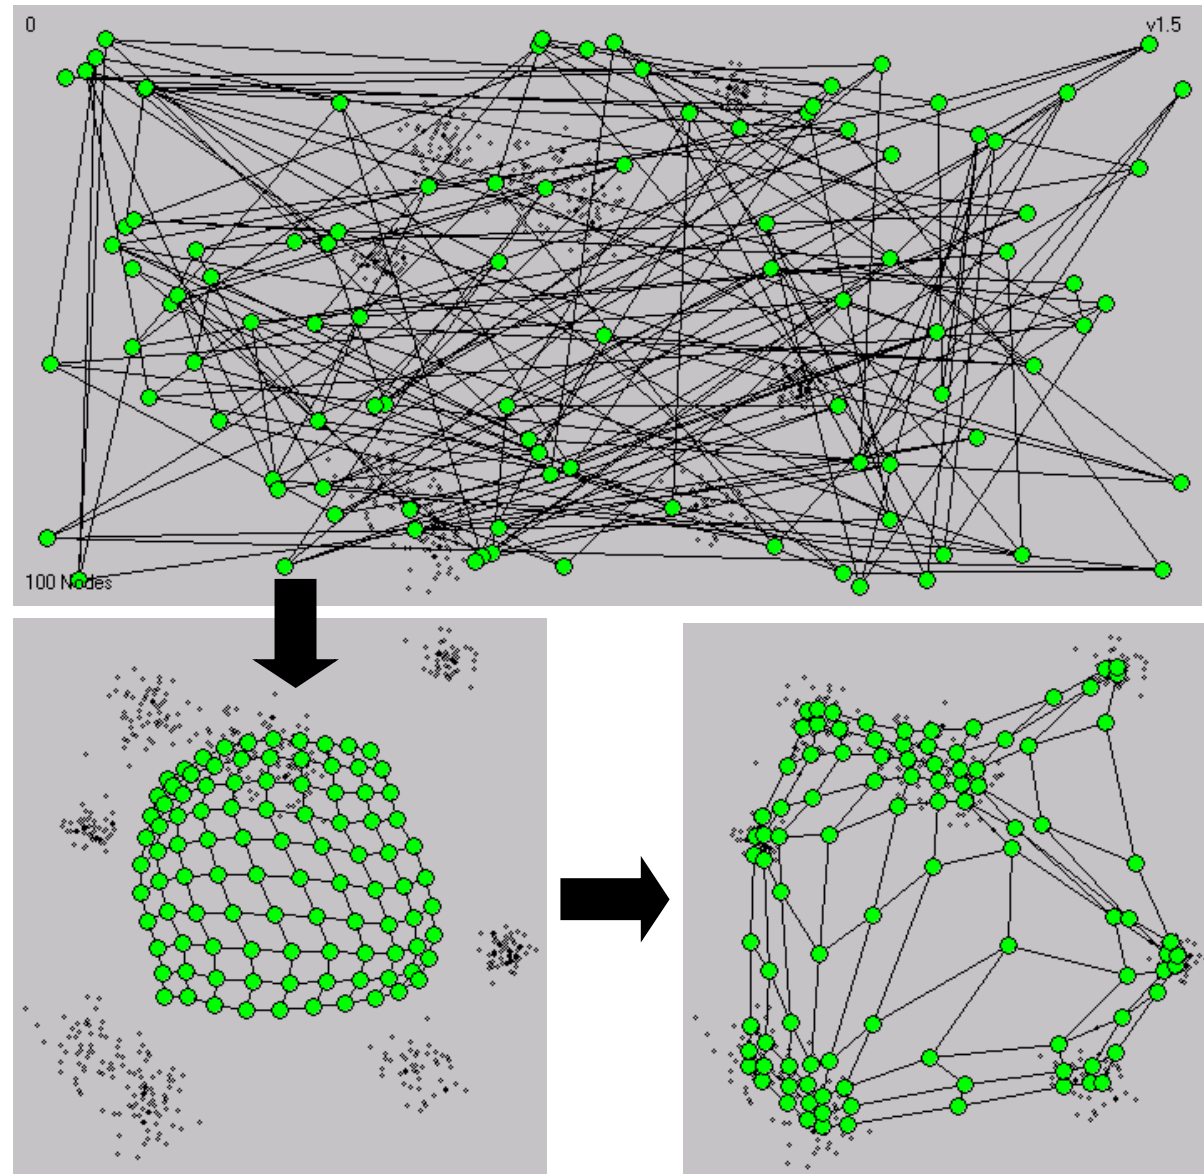
Work better for non-stationary distributions

- SOFM will follow a moving distribution, or even a jumping one
 - since winning nodes drag their neighbours along with them
- But, over time, it will gradually lose this ability!
 - due to the decaying parameters
 - (learning rate and neighbourhood function widths)
- Decaying parameters → we impose a time limit
 - for how long do we want to be plastic? (able to adapt)
- *Can we avoid decaying parameters?*




Self-Organizing Feature Maps

- It's actually OK to randomize initial positions
 - Winners will drag their neighbours along with them
- Nodes close to each other, in the map, will very quickly also get similar weight vectors
 - forming a 'fishing net', which then stretches out, trying to fit the data



Self-Organizing Feature Maps

- In SOFM, the *neighbourhood graph* is fixed (a grid)
- In this example, the 'fishing net' is very stretched in the middle
 - some under-utilized nodes there,
 - prevented from moving into a data cluster, by other neighbours pulling in other directions
- There will always be nodes which come in the middle of such conflicts
- Should they still be neighbours?
- Can we make the neighbourhood graph dynamic?
 - Create and cut edges as needed
 - like rubber bands, which snap if too stretched, or too old?



Growing Neural Gas

Fritzke, 1995


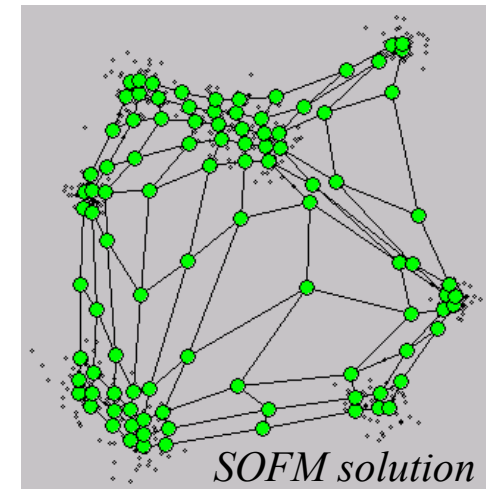
- Unsupervised growing algorithm for clustering
- Dynamic size – a growing/shrinking algorithm
 - though it grows much faster than it shrinks
- Dynamic neighbourhood – who is neighbour to whom is not fixed
 - Defined by a graph (as the grid lines in SOFM above)
- All parameters are constants!
 - GNG will not 'freeze' after a while, as SOFM does
 - Great for on-line learning and moving targets!
 - GNG will even follow jumping targets

Growing Neural Gas

Implementation considerations

We must decide ...

- Which node(s) to move, given an input
 - and by how much
- How to define and update the neighbourhood graph
 - which is no longer a fixed grid
- How to grow
 - when and where to insert new nodes



Growing Neural Gas

Implementation considerations

Which node(s) to move, given an input

- Move not only the winner (k), but also its (current) neighbours
- The winner should move with a much greater gain factor than its neighbours:

$$\begin{aligned}\Delta \bar{w}_k &= \epsilon_k (\bar{x} - \bar{w}_k) \\ \Delta \bar{w}_n &= \epsilon_n (\bar{x} - \bar{w}_n)\end{aligned}\quad \epsilon_k \gg \epsilon_n$$

- where k is the winner, n is a neighbour, \bar{x} is the input vector and \bar{w} is a node position (its weight vector)
- Both gain factors, ϵ_k and ϵ_n , are constants
- All neighbours use the same step length
(n in ϵ_n is a name, not an index)

Growing Neural Gas

Implementation considerations

Neighbourhood (and removal of nodes)

- All current neighbours are connected in a graph
- Each edge in the graph has an associated *age*
- For each input vector
 - find the closest node (k , the winner), and the second closest (r)
 - If k and r are not already connected (i.e. neighbours), connect them with a new edge
 - Set the age of the edge between k and r to 0 (zero)
 - The age of all other edges from k is incremented by one
 - If any edge becomes too old ($> a_{max}$), remove it
 - If any node loses its last edge, remove that node as well




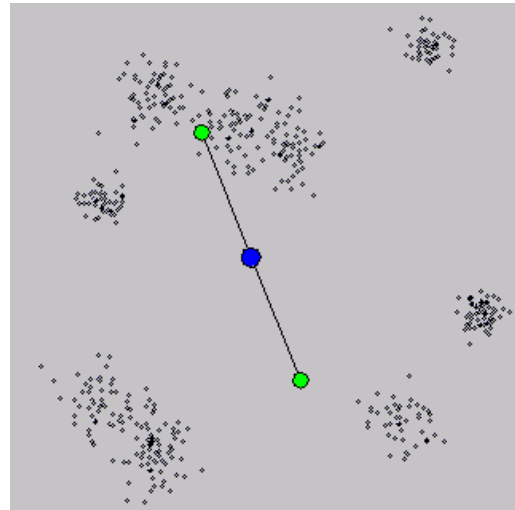

Growing Neural Gas

Implementation considerations

Growth

- Every time a winner, k , is found, add the distance (from the input) to a local error variable
 - The error is proportional to the accumulated distance this node as moved, *as a winner*
- At fixed time intervals, insert a new node where it is most likely needed:
 - halfway between the node with the largest error, and the node among its current neighbours with the largest error
- The error of a node is decayed over time
 - this is not a decaying parameter, though


GNG in action



Growing Neural Gas

Possible problem: dead units

- There is only one way for an edge to get 'younger'
 - when the two nodes it interconnects are the two closest to the input, the age is reset to 0
- If one of the two nodes wins, but the other one is not the runner-up, then, and only then, the edge ages
- If *neither* of the two nodes win, the edge does not age!



The input distribution has jumped from the lower left to the upper right corner, leaving a set of nodes in the lower left which will never be used (unless the distribution jumps back again). If they are never used, they will not be removed. Not necessarily a problem though (since it might jump back again).


Growing Neural Gas

Possible problem: dead units

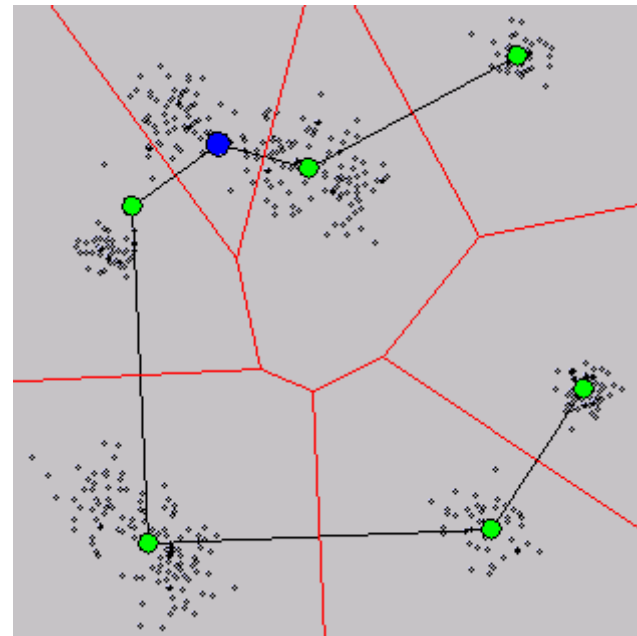
- Common extension: GNG-U
 - Removes nodes with low *utility*, based on frequency of winning and closeness to other nodes
- GNG forms a subset of a *Delaunay triangulation*
- GNG is used mostly for modelling distributions in image analysis
- Can also be used to train the hidden layer in RBF networks (lecture 15)
 - ➔ Automatic sizing of that hidden layer
- See the algorithm in Fritzkes paper for details

Delaunay triangulation

Connect the codebook vectors in all adjacent Voronoi regions



*Voronoi regions (red) and
Delaunay triangulation (yellow)*




*The neighbourhood graph in GNG is
a subset of the Delaunay triangulation*

SOFM in Matlab

Lab 3

- The grid is hexagonal, not square
- Epoch learning, instead of pattern learning
- The neighbourhood function is a tophat function!



- Step length, η , is 1!
 - This means that all neighbours in a radius around the winner, are moved to the input (not just towards it)!
 - At first glance, this should not work!
 - But in combination with the use of epoch learning, it does! (since weight changes are accumulated over the whole training set, before actually applied)