

Swarm Intelligence


Bacteria, ants,
birds and fish
in computing



Olle Gällmo




Swarm Intelligence



- Intelligence (artificial and natural) is often considered a property of *individuals*
- Most intelligent animal species are social
- Are we social because we are intelligent or is it the other way around?
 - Both, of course (they are co-evolving), but one direction has been studied longer than the other



Swarm Intelligence



- Intelligence can emerge from social interaction
- *Emergent behaviour* – when a group behaves in ways that were not “programmed” into its members
- Swarm intelligence
 - ✱ simulated social interaction
 - ✱ emergent collective intelligence in groups of simple agents

Observations

- Bird flocks and fish schools move in a coordinated way, but there is no coordinator (leader)
 - ✱ So, what decides the behaviour of a leader-less flock?
- Ants and termites quickly find a short path between the nest and a food source
 - ✱ ... and solve many other advanced problems as well
 - keeping cattle, building (ventilated) housing, coordinated heavy transports, tactical warfare, cleaning house, etc.
 - ✱ A single ant is essentially a blind, memory-less, random walker!
- Distributed systems without central control
- Useful not only to simulate but also to solve optimization problems

Bird flocks and fish schools

- Local interaction
- No leader
- Simple local rules – a weighted combination of several goals
 - ✱ match velocity of neighbours
 - ✱ avoid collisions with neighbours
 - ✱ avoid getting too far from neighbours
 - or strive for centre of the flock (fish)
- Sufficient to make very realistic simulations
 - ✱ used in movies and computer graphics
 - ✱ remove the match-velocity rule: insect swarm
 - ✱ remove collision rule: cultural interaction





UPPSALA
UNIVERSITET

Stampede in "Lion King"






What about the ants?


- How do they search for the shortest route?
 - ✱ Individual ants don't
 - ✱ The colony does
- Ant colonies are much more intelligent than ants
 - ✱ Ant colonies adapt, ants don't (much)
 - ✱ Ants have almost no memory and can not build cognitive maps. Ant colonies can (and do)
 - Mammals build cognitive maps in their brains
 - Ant colonies build them in their environment, through [pheromone trails](#)
- Ants are better thought of as cells in a greater organism – the colony
 - ✱ Also without leader – the queen is not a controller




Ants find shortest paths



At first, the ants
select at random



After a while, pheromones
become more concentrated
on the shortest route



Drawings by Marco Dorigo

Stigmergy

Indirect communication and coordination, by local modification and sensing of the environment



Walls, tunnels and bridges



Computational tools

- Cellular automata
 - ✱ 1940's (von Neumann et al)
 - ✱ an alternative computer architecture
- Ant Colony Optimization
 - ✱ 1991 (Dorigo)
 - ✱ mostly for combinatorial optimization
- Particle Swarm Optimization
 - ✱ 1995 (Kennedy & Eberhart)
 - ✱ more general optimization technique

Computational tools


- Cellular automata
 - ✱ 1940's (von Neumann et al)
 - ✱ an alternative computer architecture
- Ant Colony Optimization
 - ✱ 1991 (Dorigo)
 - ✱ mostly for combinatorial optimization
- Particle Swarm Optimization
 - ✱ 1995 (Kennedy & Eberhart)
 - ✱ more general optimization technique

Cellular automata

- Massively parallel system of identical communicating state machines (cells)
- A cell's *state* (e.g. on/off) is a function of the states of the cells it communicates with (its neighbours)
 - ✱ The neighbourhood is usually topological
- Used to model/animate fluids (e.g. the water in *Find Nemo*), gases, bacterial growth, swaying grass, social interaction, epidemics, in ecological simulations etc.

Conway's Game of Life

- World: a 2D grid. Each square represents a cell
- States: Living or dead
- Neighbourhood: The eight surrounding cells
- Initialize with a random number of living cells
- State transition rules:
 - ✱ A living cell with <2 living neighbours dies (loneliness)
 - ✱ A living cell with >3 living neighbours dies (overcrowded)
 - ✱ A dead cell with exactly 3 living neighbours comes alive
 - ✱ All other cells keep their current state



Conway's Game of Life


■ State transition rules:

- ✱ A living cell with <2 living neighbours dies (loneliness)
- ✱ A living cell with >3 living neighbours dies (overcrowded)
- ✱ A dead cell with exactly 3 living neighbours comes alive
- ✱ All other cells keep their current state

Will die



Stationary




Conway's Game of Life


■ State transition rules:

- ✱ A living cell with <2 living neighbours dies (loneliness)
- ✱ A living cell with >3 living neighbours dies (overcrowded)
- ✱ A dead cell with exactly 3 living neighbours comes alive
- ✱ All other cells keep their current state

Cyclic:




Cyclic and mobile





UPPSALA
UNIVERSITET

Life demo






Observations

- Simple (and deterministic!) rules – complex emergent behaviour
- Few rule sets have this property
 - ✱ Why?
- GA/GP to find new rules?
- Fixed set of rules (e.g. Conways)
 - ✱ still universal
 - ✱ depends on initial cell configuration
- One-way function (cryptography)
- Natural Computation, but is it ML?



CA and maze problems



- World: 2D grid which overlays the maze
- States: Corridor or wall
- Neighbourhood: The four surrounding cells (n, e, s, w)
- Initialize cells according to the maze
- State transition rules:
 - ✱ A corridor cell with 3 or 4 neighbouring wall cells becomes a wall cell
 - ✱ A wall cell remains a wall cell
- Terminates after n generations, where n is the length of the longest blind alley
- Only cells on a route between start and goal remain corridor cells
- Can extend to include several goals, finding the shortest route, etc.






UPPSALA
UNIVERSITET

Construction and repair



Other applications

- Collaborating simple robots
 - ✱ Locomotion
 - ✱ Space probes
- Modelling
 - ✱ Water, avalanches, traffic flows, ...
- Map/level generators for games



Computational tools


- Cellular automata
 - ✱ 1940's (von Neumann et al)
 - ✱ an alternative computer architecture
- Ant Colony Optimization
 - ✱ 1991 (Dorigo)
 - ✱ mostly for combinatorial optimization
- Particle Swarm Optimization
 - ✱ 1995 (Kennedy & Eberhart)
 - ✱ more general optimization technique

Ant Colony Optimization

- Family of combinatorial optimization algorithms, based on ant behaviour
- Common benchmark: the Travelling Salesman Problem (TSP)
- Common 'real' applications
 - ✱ Scheduling and
 - ✱ Network routing (AntNet)
- Members: ACS, Ant-Q, MMAS, AS_{rank} , ...
 - ✱ most of which are extensions to Dorigo's Ant System (AS)

Traveling Salesman Problems (TSP)

- Find the shortest tour through N cities, and then back to the starting point, such that
 - ✱ each city is visited once and only once
- NP-hard
 - ✱ $(N-1)!/2$ possible tours
 - ✱ exhaustive search intractable
- Specialized algorithms exist
 - ✱ and are hard to beat



Ant System for TSP


Each ant (k)

- is placed in a randomly selected city
- remembers the partial solution found so far (initially, the start city only)
- moves stochastically from city (i) to city (j), by some transition probability

$$p_{ij}^k(t)$$

which depends on


- pheromone intensity, τ_{ij}
- local information, η_{ij} (distance)
- whether j is feasible (not already visited)





Transition probabilities

Pheromone intensity, at time t , on the path from city i to j



Local information: In TSP $\eta_{ij}=1/d_{ij}$, where d_{ij} is the distance between city i and j

$$p_{ij}^k(t) = \frac{[\tau_{ij}(t)]^\alpha * [\eta_{ij}]^\beta}{\sum_{c \in C_i^k} [\tau_{ic}(t)]^\alpha * [\eta_{ic}]^\beta}, j \in C_i^k$$

Probability, at time t , of ant k traveling from city i to city j

Set of feasible destination cities (directly reachable from city i , and not yet visited by ant k)



Effects of α and β

$$p_{ij}^k(t) = \frac{[\tau_{ij}(t)]^\alpha * [\eta_{ij}]^\beta}{\sum_{c \in C_i^k} [\tau_{ic}(t)]^\alpha * [\eta_{ic}]^\beta}, j \in C_i^k$$

- If $\alpha=0, \beta>0$
 - ✱ Pheromone information discarded, only local info used
 - ✱ Stochastic greedy search with multiple starting points
- If $\alpha>0, \beta=0$
 - ✱ No local information used, only pheromones
 - more like real ants (?)
 - ✱ May lead to premature convergence
 - all ants tend to follow the same (suboptimal) route
 - difficult to discover new shortcuts (as for real ants)

Pheromone update

When all ants have completed a tour, let each ant deposit pheromones on the paths it followed

Evaporation rate
($0 < \rho \leq 1$)

Sum over all ants, k

$$\tau_{ij}(t+1) = (1 - \rho)\tau_{ij}(t) + \sum_k \Delta\tau_{ij}^k(t) \quad \forall(i, j)$$




ant k 's contribution

$$\Delta\tau_{ij}^k(t) = \begin{cases} 1 / L_k(t) & \text{if path } ij \text{ was used by ant } k \\ 0 & \text{otherwise} \end{cases}$$

$L_k(t)$ = length of ant k 's tour


Trivial example (4 cities)

Two ants



$$L_2 = 2 + 2 * \sqrt{2} \approx 4.8$$

$$\Delta\tau_{ij}^k(t) = 1/L_k(t) = 1/4.8 \approx 0.21$$



Ant System TSP Demo

- 20 cities ($19!/2 = 6.1 \cdot 10^{16}$ possible tours)
- 20 ants (one in each city)
- $\alpha = \beta = 1$
- Evaporation rate, $\rho = 0.9$



Notes on Ant Colony Opt.

- Not really a swarm?
 - ✱ These ants are not aware of each other, only of pheromones and other local info
- No direct communication \Rightarrow very scalable!
- The TSP solution demonstrated here works, but is not state-of-the-art
 - ✱ Best ACO algorithms exploit available global information
- ACO is most promising for non-stationary problems (e.g. network routing)
 - ✱ fewer competitors

What is "optimal"?

- Specialized algorithms v.s. general "black-box" ones
- Problem oriented def. of "optimal"
 - ✱ Specialized algorithms usually wins
- In practice (in industry), 'optimality' involves other concerns as well
 - ✱ time and cost to setup and maintain
 - ✱ amount of knowledge required
 - ✱ good enough is good enough

Pragmatic advice

Your algorithms teacher would probably not agree, but

It is often better to use an algorithm/method you know well, than to search for (and tune) the "best" one!

But, of course, if you happen to know the best one ...

Computational tools

■ Cellular automata

- ✱ 1940's (von Neumann et al)
- ✱ an alternative computer architecture

■ Ant Colony Optimization


- ✱ 1991 (Dorigo)
- ✱ mostly for combinatorial optimization

■ Particle Swarm Optimization

- ✱ Kennedy & Eberhart

- *Swarm Intelligence*, Morgan Kaufmann, 1995

- ✱ a more general optimization technique






Particle Swarm Optimization

- Originally intended to simulate bird flocks and to model social interaction
 - ✱ but stands on its own as an optimization tool
- A population of *particles*
 - ✱ Population sizes, typically 10-50 (smaller than in EC)
- A particle, i , has a position, \mathbf{x}_i , and a velocity, \mathbf{v}_i
 - ✱ Both vectors in n -dimensional space
- Each particle's position, \mathbf{x}_i , represents one solution to the problem
- Each particle remembers the best position it has found, so far, \mathbf{p}_i

The flying particle

- The particles "fly" through n -dimensional space, in search for the best solution

$$x_{i,d}(t) = x_{i,d}(t-1) + v_{i,d}(t)$$



- The velocities, v , depend on previous experience of this particle and that of its neighbours

Discrete time \Rightarrow velocity = step length

- Neighbourhood definition varies
 - ✱ Extreme cases: pbest (personal) and gbest (global)
 - ✱ General case: lbest (local best)
 - pbest and gbest are special cases

Personal best (pbest)

- No interaction between particles


For all particles, i , and all dimensions, d :

$$v_{i,d}(t) = v_{i,d}(t-1) + U(0, \varphi) * (p_{i,d} - x_{i,d}(t-1))$$

Uniformly
distributed
random
number

acceleration
constant
(typically ≈ 2)

best position found
so far by particle i




Neighbourhood structure

*No interaction? That's not
really a swarm, is it?*

A pbest particle in action

$$v_d(t) = v_d(t-1) + U(0, \varphi) * (p_d - x_d(t-1)) \quad \forall d$$

$$x_d(t) = x_d(t-1) + v_d(t)$$



Global best (gbest)


■ Global interaction

For all particles, i , and all dimensions, d :

$$v_{i,d}(t) = v_{i,d}(t-1) +$$

$$\underbrace{U(0, \varphi_1) * (p_{i,d} - x_{i,d}(t-1))}_{\text{Cognitive component}} + \underbrace{U(0, \varphi_2) * (p_{g,d} - x_{i,d}(t-1))}_{\text{Social component}}$$

Best solution
found so far
by *any* particle




Star neighbour-
hood structure

*To (immediately) know the global best is not very realistic, is it?
(Yes, I know the Borg do)*

A gbest particle in action

$$v_{i,d}(t) = v_{i,d}(t-1) + \underbrace{U(0, \varphi_1) * (p_{i,d} - x_{i,d}(t-1))}_{\text{possibility of } x(t), \text{ if } \varphi_1 = \varphi_2 = 2 \text{ and random numbers are drawn per element}} + \underbrace{U(0, \varphi_2) * (p_{g,d} - x_{i,d}(t-1))}_{\text{possibility if random numbers are drawn per vector}}$$




Local best (lbest)

■ Local interaction

For all particles, i , and all dimensions, d :

$$v_{i,d}(t) = v_{i,d}(t-1) + \\ U(0, \varphi_1) * (p_{i,d} - x_{i,d}(t-1)) + U(0, \varphi_2) * (p_{l,d} - x_{i,d}(t-1))$$

Best solution found so far by any particle *among i 's neighbours* (in some structure)




Ring neighbourhood structure

Nice, local, realistic, slower than gbest. Less risk of premature convergence


Simulation in 2D

Lbest with $\varphi_1=1.8$, $\varphi_2=2.3$

- Nhood: the 2 immediate neighbours
- $V_{max} = \text{range}/25$



Bowl



Alpine 2D

Observations

- Usually requires a speed limit (V_{max})
- Actual velocity (v) usually close to V_{max}
- Discrete time \rightarrow velocity = step length
- Low accuracy close to global optimum
- Decaying V_{max} ?
 - ✱ Imposes a time limit to reach the goal
- Inertia (meta-inertia, really)
- Constriction

Constriction

- Constrict swarm with a factor K
 - ✱ to avoid divergence ("explosion")
 - ✱ no longer need a speed limit
 - ✱ lowers speed around global optimum

$$v_{i,d}(t) = K * (v_{i,d}(t-1) + U(0, \varphi_1) * (p_{i,d} - x_{i,d}(t-1)) + U(0, \varphi_2) * (p_{l,d} - x_{i,d}(t-1)))$$


$$K = \frac{2}{\varphi - 2 + \sqrt{\varphi^2 - 4\varphi}}, \text{ where } \varphi = \varphi_1 + \varphi_2 > 4$$

- K is a function of φ_1 and φ_2 only $\Rightarrow K$ is constant
 - ✱ yet it gives the swarm a converging behaviour, as if K was a decaying variable

Simulation with K

Lbest with $\varphi_1=1.8$, $\varphi_2=2.3 \Rightarrow K=0.7298$

- Nhood: the 2 immediate neighbours
- No V_{max}



Binary particle swarms

- Velocities updated as before
- The positions:

$$x_{i,d} = \begin{cases} 1, & \text{if } U(0,1) < \text{logistic}(v_{i,d}) \\ 0, & \text{if } U(0,1) \geq \text{logistic}(v_{i,d}) \end{cases} \quad \text{logistic}(v_{id}) = \frac{1}{1 + e^{-v_{i,d}}}$$

- $V_{max} = \pm 4$
 - ✱ in order not to saturate the sigmoid
 - ✱ so that there is at least some probability (0.018) of a bit flipping

What makes PSO special?

- Its simplicity
- Adaptation operates on velocities
 - ✱ Most other methods operate on positions
 - ✱ Effect: PSO has a builtin momentum
 - ✱ Particles tend to hurdle past optima – an advantage, since the best positions are remembered anyway
- Few parameters to set, stable defaults
- Relatively easy to adapt swarm size
- Not very good for fine-tuning, though constriction helps

Notes on PSO

- Many publications are misleading on one important point:
 - ✱ The random numbers should be drawn *per element* (not per vector)
- Not really a swarm (though it behaves as one if V_{max} is small)
 - ✱ Particles don't know other particles positions or velocities, only their personal bests
- The neighbourhood in *lbest* is structural (social)
 - ✱ Could it be topological (geographical)?


PSO and neural networks

- PSO can be used to train neural networks
- Each particle represents one network
 - ✱ x = a vector of all networks weights
 - ✱ also node type can be parameterized
- Hypothesis: This works better than EC
 - ✱ Crossover in EC may not be sound (for this)
 - Two individuals selected for crossover may have very little in common
 - ✱ Mutation in EC corresponds to moving the individual, but it is not directed, as in PSO
 - Both are random, but distribution depends on the situation in PSO, not so (usually) in EC
 - ✱ PSO population typically smaller than in EC
 - Speed advantage, since evaluation is a bottle-neck



UPPSALA
UNIVERSITET

MLP+PSO+XOR



Grammatical Swarm

■ Grammatical Evolution (GE) ...

BNF Grammar


$\langle \text{expr} \rangle ::= \langle \text{const} \rangle \mid \langle \text{expr} \rangle \langle \text{op} \rangle \langle \text{expr} \rangle$

$\langle \text{const} \rangle ::= 0 \mid 1 \mid 2$

$\langle \text{op} \rangle ::= + \mid -$

Genome

23 | 4 | 17 | 42 | 62 | 10 | 6 | 22 | ...



... where the integer string is trained by PSO instead of a genetic algorithm (GA)

Room for thought ...

- High-dimensional spaces
- Bias along coordinate system axis
- Multiplying by vector instead of by element – effects in practice
- "PSO on Ice" (PSO-GD)
- Dynamic neighbourhood, like GNG
- Parameter free PSO (Tribes)
- Multiple swarms and multiple goals



UPPSALA
UNIVERSITET

"Once again, nature has provided us with a technique for processing information that is at once elegant and versatile"

/ Kennedy & Eberhart -95